

Shivam Singh

✉ shivamsingh3458@gmail.com | [in LinkedIn](#) | [GitHub](#) | shivhere.github.io | [Google Scholar](#) | Phone: +91 8130870896

EDUCATION

M.S. -- Electronics and Communication Engineering | IIIT HYDERABAD (2023–2026)

CGPA: 8.33/10

B.Tech -- Electronics and Communication Engineering | RKGIT, GHAZIABAD (2017–2021)

CGPA: 8.07/10

RESEARCH & PROFESSIONAL EXPERIENCE

AMD INDIA PRIVATE LIMITED | RESEARCH FELLOW

AI Models & Agents Team | Jun 2025 - July 2026

Supervisors: Saptarshi Majumder, Pratik Prabhanjan Brahma

- **Pause and Think: Video-Grounded Assistive Action Suggestion:** 2025
 - Proposed a reasoning-centric framework for video-grounded assistive action suggestion by introducing *Pause-and-Think-T* (~10k structured reasoning QA pairs) with explicit reasoning supervision and *Pause-and-Think-B*, a **300-sample** benchmark for grounded contextual understanding and goal-oriented planning across Epic-Kitchens, Assembly101, and Ego4D.
 - Designed a scalable data generation pipeline comprising annotation refinement with GPT-OSS-120B, goal-oriented clip segmentation (5–20s), and structured `<think>/<answer>` QA generation using Qwen3-VL-235B with majority voting and self-consistency verification.
 - Fine-tuned a compact **4B** transformer-based VLM using LLaMA-Factory with DeepSpeed on 8×AMD Instinct MI325 GPUs and served with vLLM, achieving **58.0%** accuracy—matching Qwen3-VL-235B (**58.9%**, **59× larger**) and GPT-5.2 on scene understanding, outperforming GPT-4o, and demonstrating strong zero-shot generalization on EgoThink and TempCompass without benchmark-specific training.
 - Deployed the optimized model on an **AMD Ryzen™ AI** platform for real-time on-device inference, demonstrating efficient edge deployment of reasoning-capable Vision-Language Models.
 - **First Author:** Led the research, algorithm design, dataset creation, model training, evaluation, deployment, and manuscript preparation. [Code ↗ | Paper ↗]
- **Spatial Video Understanding from Raw RGB Videos:** 2026
 - Proposed *SlotMem*, a memory-augmented transformer framework for object tracking, spatial understanding, and long-horizon video reasoning from raw monocular RGB videos.
 - Developed a scalable pipeline combining Depth-Anything-3, YOLO-World, and Vision-Language Models to reconstruct 3D scenes, estimate camera trajectories, and generate supervision for spatial, temporal, and goal-oriented reasoning.
 - Trained SlotMem using our automatically generated supervision together with diverse videos from VLM-3R and VICA-322K, achieving **state-of-the-art** performance on OST-Bench (**44.4%**) and outperforming all prior baselines.

INTERNATIONAL INSTITUTE OF INFORMATION TECHNOLOGY - HYDERABAD | M.S. RESEARCHER

Robotics Research Centre, Hyderabad | Jan 2023 - May 2025

Direct Supervisor: Prof. Madhava Krishna

Collaborators: Dr. Mohan Sridharan, Dr. Krishna Murthy Jatavallabhula, Dr. Brojeshwar Bhowmick

- **Task Anticipation using LLMs (IEEE ICRA 2024)** 2023
 - Co-author & Project Lead:** Developed an intelligent household agent that performs long-horizon planning by anticipating user actions through learning task execution patterns from user behavior and preferences, leveraging Large Language Models (LLMs) and Planning Domain Definition Language (PDDL)-based task and action representations. [Website ↗ | Paper ↗]
- **Task Anticipation for Multi-Agent Systems (ICRA & RSS 2024 Workshops)** 2024
 - Project Lead:** Extended LLM-based task anticipation to long-horizon planning in multi-agent settings, enabling robotic agents to anticipate domain-specific tasks, adapt to unexpected changes in human behavior, and handle environmental uncertainties for effective human-robot collaboration. [Website ↗ | Paper ↗]

- **LLM & Knowledge Graphs for Task Decomposition** (*IEEE ICRA 2025*) 2025
Project Lead: Designed a framework that integrates Large Language Models (LLMs), Knowledge Graphs (KGs), and Human-in-the-Loop (HITL) feedback to enable embodied agents to perform long-horizon planning for unseen tasks and progressively adapt to new environments. [Website [↗](#) | Paper [↗](#)]
- **Hybrid Planning with LLMs & Probabilistic Models** (*ECMR 2025*) 2025
Mentor: Guided framework design and ablation experiments for task planning under uncertainty with probabilistic effects in multi-agent settings. [Paper [↗](#)]

ALSOENERGY – INDIA | TECHNICAL SUPPORT ENGINEER

April 2021 | Dec 2021

- Delivered expert troubleshooting and step-by-step diagnostic support for clean energy management platforms (e.g., PowerTrack), collaborating with commercial solar clients to efficiently resolve complex issues regarding system performance, data visualization, and third-party hardware integration.

PUBLICATIONS & WORKSHOPS

- **Anticipate & Act: Integrating LLMs and Classical Planning for Efficient Task Execution in Household Environments:** [↗](#) R Arora*, **Shivam Singh***, K Swaminathan, A Datta, S Banerjee, B Bhowmick, KM Jatavallabhula, M Sridharan, M Krishna.
Accepted at IEEE International Conference on Robotics and Automation (ICRA) 2024 [↗](#), Yokohama, Japan.
- **AdaptBot: Combining LLM with Knowledge Graphs and Human Input for Generic-to-Specific Task Decomposition and Knowledge Refinement:** [↗](#) **Shivam Singh**, K Swaminathan, N Dash, R Singh, S Banerjee, M Sridharan, M Krishna.
Accepted at IEEE International Conference on Robotics and Automation (ICRA) 2025 [↗](#), Atlanta, US.
- **Anticipate, Adapt, Act: A Hybrid Framework for Task Planning:** [↗](#) N Dash, A Kaura, **Shivam Singh**, R Singh, S Banerjee, M Sridharan, M Krishna.
Accepted at IEEE European Conference on Mobile Robots (ECMR) 2025.
- **Pause and Think: A Dataset and Benchmark for Video-Grounded Assistive Action Suggestion:** [↗](#) **Shivam Singh**, S Majumder, P Prabhanjan, Z Liu, E Barsoum.
Accepted at IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS) 2026 [↗](#). Work done at AMD (Advanced Micro Devices, Inc.).
- **DaTAPlan: Data-driven Task Anticipation and Knowledge-driven Planning for Human-robot Collaboration:** [↗](#) **Shivam Singh***, K Swaminathan*, R Arora*, R Singh, A Datta, D Das, S Banerjee, M Sridharan, and M Krishna.
Accepted in ICRA 2024 Workshop on Cooking Robotics: Perception and motion planning [↗](#), Yokohama, Japan.
- **Anticipate & Collab: Data-driven Task Anticipation and Knowledge-driven Planning for Human-robot Collaboration** [↗](#): **Shivam Singh***, K Swaminathan*, R Arora*, D Das, S Banerjee, M Sridharan, M Krishna. *Accepted in RSS2024 Workshop on Task Specification for General-Purpose Intelligent Robots [↗](#), Delft, Netherlands*

PATENTS

- **Systems and Methods for Data-driven Task Anticipation and Knowledge-driven Planning for Human-Robot Collaboration** [↗](#): **Shivam Singh**, K Swaminathan, R Arora, R Singh, A Datta, D Das, S Banerjee, M Sridharan, and M Krishna.
US Patent App. 19/089,689, published November 6, 2025.

TECHNICAL SKILLS

Research Areas: Foundation Models, Large Language Models (LLMs), Vision-Language Models (VLMs), Reasoning Models, Multimodal Learning, Agentic AI, Temporal Reasoning, Memory-Augmented Models, Embodied AI

Model Training & Inference: PyTorch, Hugging Face Transformers, Hugging Face Datasets, LLaMA-Factory, SFT, PEFT/LoRA, DeepSpeed, Distributed Training, vLLM, Tensor Parallelism, Efficient Inference, ROCm

Data, Evaluation & Infrastructure: Multimodal Data Curation, Synthetic Data Generation, Dataset Construction, Benchmarking, Model Evaluation, Docker, Slurm, OpenCV, Open3D

Programming & Robotics: Python, C++, C, PDDL, RDDL, ROS, Gazebo, AI2-THOR